

NARANI HEMANTH KUMAR

+91 93982 42096 · hemanthkumar0330@gmail.com · Hyderabad, India · [linkedin.com/in/narani-hemanth-kumar](https://www.linkedin.com/in/narani-hemanth-kumar)

Professional Summary

Results-driven Data Engineer with 4 years of hands-on experience designing, building, and optimising enterprise-scale data pipelines, ETL workflows, and cloud-based data platforms on Microsoft Azure. Delivered end-to-end data engineering solutions across Microsoft Store, M365, Xbox, Windows, and Office data domains — processing 5TB+ of daily commerce transactions with 95%+ SLA compliance. Deep expertise in PySpark, Azure Data Factory, Azure Databricks, Delta Lake, Medallion Architecture, and Azure Synapse Analytics, with a proven track record of reducing pipeline execution times by 40%+, automating incident recovery workflows (reducing manual effort by ~60%), and building Power BI dashboards consumed by senior stakeholders across Microsoft's digital commerce division.

Core Competencies

✓ Azure Data Factory (ADF)	✓ Azure Databricks & PySpark	✓ Azure Synapse Analytics
✓ Azure Data Lake Storage (ADLS)	✓ Delta Lake & Medallion Architecture	✓ Spark Structured Streaming
✓ ETL / ELT Pipeline Development	✓ Data Warehousing & Data Modeling	✓ Star & Snowflake Schema
✓ SQL & MS SQL Server	✓ Data Quality & Data Governance	✓ Data Lineage & Schema Enforcement
✓ Python & PySpark	✓ CI/CD via Azure DevOps	✓ Power Automate & Workflow Automation
✓ Power BI (DAX, Power Query)	✓ Root Cause Analysis (RCA)	✓ ServiceNow & IcM Incident Management
✓ Azure Monitor & Log Analytics (KQL)	✓ Pipeline Monitoring & SLA Management	✓ Agile Delivery & Scrum

Work Experience

Data Engineer December 2021 – November 2025

Tata Consultancy Services · Client: Microsoft Corporation

Project: Microsoft Digital Stores Support | Domains: Microsoft Store · Office · M365 · Windows · Gaming / Xbox

Data Pipeline Architecture & ETL Development

- Designed and deployed 20+ high-throughput Azure Data Factory (ADF) pipelines ingesting 5TB+ of daily commerce data from REST APIs, RDBMS, and flat files into Azure Data Lake Storage Gen2 (ADLS Gen2), serving as the single source of truth across five Microsoft product domains.
- Architected a three-tier Medallion Architecture (Bronze → Silver → Gold) in Azure Databricks using PySpark, enforcing schema validation, data quality checks, and schema evolution rules — reducing data defect rate by 35% and improving data reliability for downstream consumers.
- Built scalable ELT workflows in Azure Synapse Analytics using T-SQL and Synapse Pipelines to load structured data into a Star Schema model, enabling reliable, low-latency reporting for business intelligence teams.
- Implemented Spark Structured Streaming on Azure Databricks with Delta Lake to enable near-real-time data ingestion for high-velocity commerce events, maintaining ACID transactions, data lineage tracking, and full auditability across ADLS Gen2 and Azure SQL.
- Managed structured data storage across ADLS Gen2, Delta Lake, and Azure SQL — applying partitioning strategies, Z-ordering, and data compaction to optimise read/write performance for business-critical operations.

Performance Optimisation & Cost Efficiency

- ▶ Reduced daily pipeline execution time by 40%+ by rewriting complex PySpark jobs with broadcast joins, partition pruning, adaptive query execution (AQE), and intelligent caching — cutting Azure Databricks compute costs and significantly improving SLA adherence for sales and inventory data.
- ▶ Optimised 50+ complex SQL queries and Synapse Analytics scripts, reducing query execution time by up to 55% through indexing, statistics updates, and query plan analysis.
- ▶ Performed regular Azure Cost Management reviews on Databricks clusters and ADF pipeline runs, identifying and eliminating over-provisioned resources — contributing to measurable reductions in monthly cloud spend.

Monitoring, Reliability & Production Support

- ▶ Maintained 95%+ SLA compliance across all production pipelines by automating alerting, failure detection, and process-recovery workflows in Microsoft Power Automate — reducing mean time to recovery (MTTR) and cutting manual toil by ~60%.
- ▶ Resolved 200+ production incidents by performing Root Cause Analysis (RCA) on data pipeline failures via ServiceNow and IcM — documenting post-mortems, implementing preventive controls, and minimising impact on Microsoft Store's live commerce operations.
- ▶ Monitored enterprise pipeline health and job executions using Azure Monitor, Log Analytics (KQL queries), and Geneva Health Portal — proactively identifying bottlenecks and alerting engineering teams before SLA breaches occurred.
- ▶ Provided 24/7 on-call production support for Tier-1 data pipelines, managing incident workflows end-to-end from detection through root cause resolution and stakeholder communication.

DevOps, Automation & Cross-Functional Collaboration

- ▶ Implemented CI/CD pipelines via Azure DevOps — configuring Git-based branching strategies, automated testing, and release gates that reduced deployment errors by ~50% and accelerated feature delivery across data engineering sprints.
- ▶ Built 10+ interactive Power BI dashboards using DAX measures and Power Query transformations, visualising product sales trends, discount effectiveness, inventory KPIs, and operational metrics — used weekly by senior Microsoft stakeholders for strategic decisions.
- ▶ Collaborated with cross-functional engineering, business intelligence, and product teams in Agile/Scrum sprints to translate business requirements into scalable data architecture — continuously iterating on data models and pipeline designs to meet evolving commerce analytics needs.

Technologies: Azure Data Factory · Azure Databricks · PySpark · Python · SQL / T-SQL · MS SQL Server · Azure Synapse Analytics · ADLS Gen2 · Delta Lake · Spark Structured Streaming · Azure Monitor · Log Analytics (KQL) · Power BI · Power Automate · Azure DevOps · ServiceNow · IcM · Geneva Health Portal · Medallion Architecture

Certifications & Professional Development

Complete Data Analytics with Azure, AI, and Copilot — Basic to Advanced

October 2024

Grow Data Skills

Skills: SQL, Excel, Statistics, Power BI, Python, Azure Data Analytics, AI-enabled Analytics, Copilot Integration

Power BI A-Z: Hands-On Power BI Training for Data Science

Udemy

Skills: Data Visualisation, Dashboard Development, DAX, Power Query, Trend Analysis, KPI Reporting

Get Started with Power Automate + Build Approval Flows with Power Automate

Microsoft (2 Certificates)

Skills: Recurring Flows, Scheduled Flows, Event-Triggered Automation, Approval Process Workflows

Azure Data Engineering End-to-End

Online Certification

Skills: Azure Data Factory, Azure Databricks, PySpark, Delta Lake, SQL Data Warehousing, Azure DevOps, Big Data Fundamentals, Python

Generative AI for Beginners

Microsoft / Online

Skills: Prompt Engineering, Large Language Models (LLMs), AI-enabled Data Analytics, Responsible AI

Education

Bachelor of Technology — Aeronautical Engineering

2017 – 2021

Malla Reddy College of Engineering and Technology, Hyderabad | CGPA: 8.56 / 10

Board of Intermediate — Mathematics, Physics & Science

2015 – 2017

Sri Chaitanya Junior Kalasala | 89%

Secondary School Certificate (Grade 10)

2015

Newgen School of Excellence | CGPA: 8.8 / 10